

# Accurate Pouring with an Autonomous Robot Using an RGB-D Camera

Chau Do and Wolfram Burgard

University of Freiburg, Autonomous Intelligent Systems, Georges-Köhler-Allee 80,  
79110 Freiburg, Germany, {do,burgard}@informatik.uni-freiburg.de

**Abstract.** Robotic assistants in a home environment are expected to perform various complex tasks for their users. One particularly challenging task is pouring drinks into cups, which for successful completion, requires the detection and tracking of the liquid level during a pour to determine when to stop. In this paper, we present a novel approach to autonomous pouring that tracks the liquid level using an RGB-D camera and adapts the rate of pouring based on the liquid level feedback. We thoroughly evaluate our system on various types of liquids and under different conditions, conducting over 250 pours with a PR2 robot. The results demonstrate that our approach is able to pour liquids to a target height with an accuracy of a few millimeters.

**Keywords:** liquid perception, robot pouring, household robotics

## 1 Introduction

A capable and effective domestic service robot must be able to handle everyday tasks involving liquids. Some examples are filling a cup or measuring out a certain amount of liquid for baking or cooking. This requires the ability to perceive the liquid level while pouring and using this information to decide when to stop pouring. This is a challenging task, considering the large selection of liquids available and their varying characteristics. In this paper, we consider the problem of tracking a liquid and determining when to stop pouring using depth data from a low-cost, widely available RGB-D camera.

Visually, liquids change their appearance depending on the environment, which makes it particularly challenging to detect the fill level of a container using vision. In this paper we therefore investigate the estimation of the fill level based on data provided by a depth camera such as the ASUS Xtion Pro or Microsoft Kinect. However, even for this modality the task at hand is complicated by the fact that liquids have different appearances depending on their transparency and index of refraction. For transparent liquids such as water and olive oil, the infrared light is *refracted* at the liquid boundary, causing the resulting liquid level to appear lower than it actually is. In the case of water, depending on the view angle, a full cup appears one-third to one-half full based on the depth data. On the other hand, opaque liquids such as milk and orange juice,



(a) PR2 with cups and bottles used in the experiments      (b) PR2 pouring water

**Fig. 1.** PR2 robot with bottles and cups used in the experiments (a) and close-up of a water pouring trial (b).

reflect the infrared light and the resulting depth measurement correctly represents the real liquid height. The approach presented in this paper is designed to deal with opaque and transparent liquids, by switching between the detection approaches depending on the type of liquid. We utilize a Kalman filter handling the uncertainty in the measurements and for tracking the liquid heights during a pour. For controlling the pour, we employ a variant of a PID controller that takes the perception feedback into account. We demonstrate the effectiveness of our approach through extensive experiments, which we conducted using a PR2 robot depicted in Fig. 1.

The two main contributions of this paper are a novel approach for pouring liquids into a container up to a user-defined height and the extensive analysis of the approach with respect to a large variety of parameters of the overall problem.

## 2 Related Work

A substantial amount of research has been carried out on various aspects of pouring. For example, Pan and Manocha [1] and Tamosiunaite et al. [2] focus on learning a pouring trajectory, but do not consider the problem of pouring to a specific height. In the area of learning from humans, Langsfeld et al. [3] and Rozo et al. [4] demonstrate how to pour a specific volume, and learn the pouring motion parameters for that volume. However, in their work no perception of the liquid is used. Regarding the area of liquid perception, Elbrechter et al. [5] focus on the problem of detecting liquid viscosity. Morris and Kutulakos [6] look at reconstructing a refractive surface, but this requires a pattern placed underneath the liquid surface. Mottaghi et al. [7] use deep learning to infer volume characteristics of containers. Neither of these papers consider the problem of tracking

a liquid level and pouring to a specific height. Yamaguchi et al. [8] focus on the planning side of pouring. They detect the fill level using a plastic cup, modified such that the back half is colored and use color segmentation. Our approach does not require modified containers. Yamaguchi et al. [9] use a stereo camera and apply optical flow detection to determine the liquid flow. They do not apply their approach to the problem of pouring to a specific height. Furthermore, flow detection cannot be used to determine static liquid already present in a container, which our approach is capable of as shown by our experiments. Yoshitaka et al. [10] use an RGB-D camera to detect the presence and height of liquid in a cup. For transparent liquids, they outlined a relationship between the measured height from the depth data and the real liquid height. They did not consider the problem of pouring liquids and only dealt with the detection of static liquids.

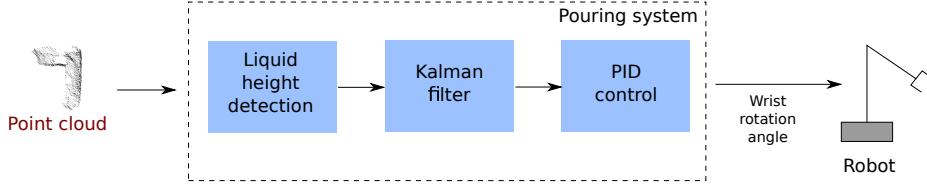
In our previous work [11], we presented a probabilistic approach to liquid height detection, which makes use of both RGB and depth data. It assumes no knowledge of the liquid type and only considers a static liquid height (i.e., non-pouring) scenario. This previous approach uses multiple images and point clouds of the static liquid, taken from different viewing angles. Capturing data from different viewing angles is required to disambiguate between liquid types and accurately determine the height, since no prior knowledge of the liquid is given. In this paper, in contrast to our prior work, we determine the height of the liquid during the pour and without observing the cup from multiple viewpoints.

Schneck and Fox [12] demonstrate it is possible to segment and track liquids using deep learning. However, this was only shown with synthetic data. In follow-up work, Schneck and Fox [13] consider the problem of pouring to a specific height. They train a deep network to classify image pixels as either liquid or not liquid. For ground truth, they synchronized a thermal camera with an RGB-D camera and recorded pouring data using water heated to 93 °C. Using this detection they then track the liquid volume while pouring. For 30 pours, they reported a mean error of 38 ml. As we will demonstrate in our experimental evaluations, our approach is able to achieve a substantially lower mean error while only employing an RGB-D camera and not requiring an expensive thermal camera as well as the heating of the liquid for collecting the training data.

Compared to these previous approaches, the method presented in this paper allows accurate pouring of liquids to given levels using a low-cost RGB-D sensor. It can be applied to different types of liquids without any need for training and is usable in every environment in which an RGB-D camera such as the Kinect or Xtion can operate.

### 3 The Autonomous Pouring System

Fig. 2 depicts our PR2 robot used for evaluating our pouring approach. We control the pouring by rotating the angle of the robot’s wrist joint. As input the system receives the point clouds from the robot’s RGB-D sensor. The cup should be positioned such that the camera can look into it and can see part of the cup bottom.



**Fig. 2.** Overview of the approach: the sensor inputs are point clouds from an RGB-D camera. The output is a rotation angle for the wrist joint of the robot to pour the liquid.

### 3.1 Liquid Height Detection

In our approach, the detection and tracking of the liquid level relies only on the depth information from an RGB-D camera such as a Kinect or an ASUS Xtion Pro. Throughout this paper, we assume that the robot knows which liquid is being poured and which target height is required. Typically, this information can easily be provided by the user, for example, by issuing a command such as “pour me a full cup of water” or following a recipe for baking and cooking.

There are two cases to consider with respect to the liquid type: opaque and transparent liquids. Opaque liquids, such as milk and orange juice, are accurately represented in the point cloud and the extracted height can be taken as the true liquid height. Transparent liquids such as water and olive oil, however, refract the light and the point cloud height is incorrect. It is possible to estimate the actual liquid height given a refracted depth measurement using a relationship based on the view angle and the index of refraction of the liquid (see Yoshitaka et al. [10] and Do et al. [11] for a full derivation). This relationship for finding the liquid height  $h$  is given by:

$$h = \left( \frac{\sqrt{n_l^2 - 1 + \cos^2(\alpha)}}{\sqrt{n_l^2 - 1 + \cos^2(\alpha)} - \cos(\alpha)} \right) h_r. \quad (1)$$

Thus, we estimate the liquid height  $h$  given the raw depth measurement height  $h_r$ , index of refraction of the liquid  $n_l$  and the angle  $\alpha$ , where  $h_r$  is the liquid height determined from the point cloud and  $\alpha$  is the incidence angle of infrared light from the camera projector with respect to the normal of the liquid surface. The latter is also determined from the point cloud. The index of refraction  $n_l$  is determined from the liquid type, which is provided by the user. It should be noted that Eq. (1) is an approximation and does not account for all physical effects.

To determine the liquid height from the depth data, we first need to detect the cup. We achieve this by finding the plane representing the table and extracting all points above the table. Then we use RANSAC [14] to determine the cylinder model for the cup. In order to avoid detecting the cup rim as part of the liquid height, we determine the diameter of the cup from the detected cylinder model and only consider a reduced diameter section for the liquid. In other words, we

only search an area defined by the cup height and the reduced diameter. Finally we extract the points inside this area and average over them to get the raw measured height.

### 3.2 Kalman Filter

To track the liquid height and filter the typically noisy depth measurements from the RGB-D camera [15], we use the Kalman filter for tracking. Thereby we assume that the liquid height follows a constant velocity motion given by

$$x_k = Fx_{k-1} + w_k, \quad (2)$$

where  $w_k \sim \mathcal{N}(0, Q_k)$  is zero-mean Gaussian noise with covariance  $Q_k$ , which represents the system noise and is given by

$$Q_k = q \begin{bmatrix} \frac{1}{3}\Delta t^3 & \frac{1}{2}\Delta t^2 \\ \frac{1}{2}\Delta t^2 & \Delta t \end{bmatrix}. \quad (3)$$

Here,  $x_k$  and  $F$  are given by

$$x_k = \begin{bmatrix} h_k \\ \dot{h}_k \end{bmatrix} \text{ and } F = \begin{bmatrix} 1 & \Delta t \\ 0 & 1 \end{bmatrix}. \quad (4)$$

The term  $h_k$  refers to the liquid height and  $\dot{h}_k$  refers to the rate of change in liquid height at time  $k$ , while  $\Delta t$  is the time interval. Note that we chose not to model a system input, as this term and its effect on the system depends on several factors such as the bottle opening, amount of liquid in the bottle and the tilt angle of the bottle. Modeling this term from recorded data would be unique for that situation only. We determined the measurement noise covariance matrix  $R_k$  and  $q$  from collected pouring data.

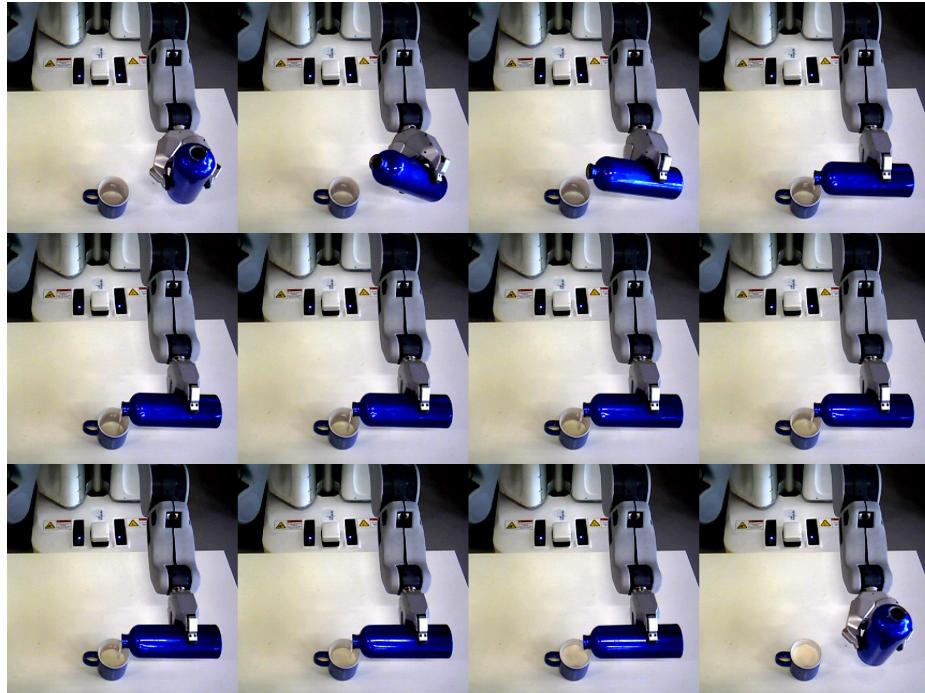
### 3.3 Pouring Control

To realize the pour, we need to determine the rotation angle of the wrist joint. In our approach we employ a variant of the PID controller that we augmented by additional control policies to ensure a smooth pour. In the event that no liquid height is detected in the point cloud, we slow down the pouring. If no liquid height update is received (i.e., the detection is too slow), we maintain the rotation angle so as not to continue pouring blindly. As soon as we detect that the required liquid height has been reached, we rotate the bottle back to its initial position.

## 4 Experiments

We implemented our approach on a PR2 robot equipped with an ASUS Xtion Pro camera and carried out a series of experiments to evaluate our approach in

real-world settings. The different cups and bottles involved in the experiments can be seen in Fig. 1. For the experiments, we positioned the cup relative to the camera with its bottom center point at an approximate horizontal distance of 25 cm and an approximate vertical distance of 75 cm. The vertical angle of the camera was chosen so that the cup bottom could be detected in the depth data. In all experiments, we placed the bottle in the gripper of the PR2 and positioned it close to the cup. Before the experiments, we fine-tuned the parameters of our PID controller. We found that proportional control was sufficient for achieving accurate results. A pouring trial for milk can be seen in Fig. 3



**Fig. 3.** Pouring sequence for a milk pouring trial (best viewed enlarged).

In total, we performed six different types of experiments, which we designed to analyze how changes in liquid type, initial liquid volume, target height, bottles and cups affect the pouring accuracy. In all of these experiments we used the same parameter values. Overall, we performed a total of 290 pours. To measure the ground truth data we hand measured the liquid height with a ruler.

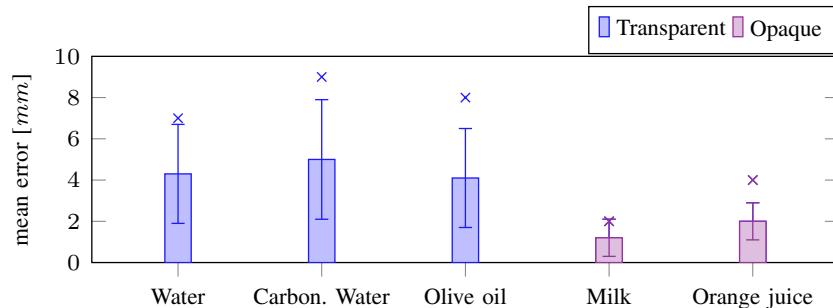
#### 4.1 Influence of Different Liquids

To analyze the impact of the liquid type, we performed experiments with water, carbonated water, olive oil, milk and orange juice. The first three are transparent

liquids, while the latter two are opaque. Overall, they represent differences in properties such as viscosity, index of refraction and carbonation.

For each liquid we poured ten times. For each pour, the initial volume in the bottle was randomly chosen from the values of [200 ml, 250 ml, 300 ml, 350 ml, 400 ml, 450 ml, 500 ml] and the target height in the cup was randomly chosen from the range of [20 mm, 30 mm, 40 mm, 50 mm, 60 mm, 70 mm]. We have the criteria that the initial volume in the bottle should be at least 100 ml more than the target height, to make sure the robot does not pour all the liquid out. To focus only on the changes caused by the liquids, we used the blue bottle and blue cup shown on the rightmost side of Fig. 1 for all the pours in this experiment.

Fig. 4 shows the absolute mean error and standard deviation for the ten pours of each liquid. The crosses show the maximum error that occurred. In almost all cases the robot over-pours. We believe this is due to the fact that we send the *STOP POUR* signal when we detect that the desired liquid height has been reached. However, this is minimized by the PID controller, which reduces the amount poured as the liquid height approaches the target height. The opaque liquid results show that this delay accounts for only around 2 mm additional liquid. Between the three transparent liquids there is not much variation and the same can be said of the two opaque liquids.



**Fig. 4.** Mean error and standard deviation for ten pours of different liquids with varying initial volume and target height. The crosses mark the maximum error.

As can be seen, the transparent liquids have higher mean errors and standard deviations versus the opaque ones. To investigate this further, we additionally poured 60 mm of water and milk into the same blue cup and recorded 200 depth measurements. The results can be seen in Table 1. In this table, “raw” refers to the raw depth measurement taken from the point cloud. In the case of water, this is before being transformed by Eq. (1). For the raw measurements, the standard deviation is about the same for both water and milk. The transformed values for water have a higher standard deviation. Looking at Eq. (1), it is clear that any noise in the raw measurement  $h_r$  is magnified by the term before it. For the recorded data, the term before  $h_r$  is around 4.03. A second observation, is that the transformed value for water underestimates the liquid height. We note

that Eq. (1) is only an approximation and does not factor in all physical effects. Accordingly, we can expect a higher mean error and standard deviation for the transparent liquids versus the opaque ones.

**Table 1.** Static Error - Target Height 60 mm

| Measurement | $\mu \pm \sigma$ [mm] |
|-------------|-----------------------|
| Water (raw) | $18.64 \pm 0.18$      |
| Water       | $58.77 \pm 0.56$      |
| Milk (raw)  | $60.37 \pm 0.12$      |

#### 4.2 Influence of Varying Initial Volume

In this experiment, we investigate how robust our system is to changes in initial volumes in the bottle. We kept the target height fixed at 40 mm and varied the initial volume by changing it to one of [350 ml, 400 ml, 450 ml, 500 ml]. In other words, for each volume we poured ten times to a height of 40 mm. As before, we only used the blue bottle and blue cup. The results for transparent liquids depicted in Fig. 4 are very similar and the same can be said of the opaque liquids. Hence for this experiment and the following ones we only used water and milk.

The top figure in Fig. 5 shows the results of this experiment for the pours. There is a slight increase in mean error as the volume increases. For water, the difference in mean error between 500 ml and 350 ml is 2.4 mm and for milk, it is 1 mm, which are not very large differences in view of the overall results.

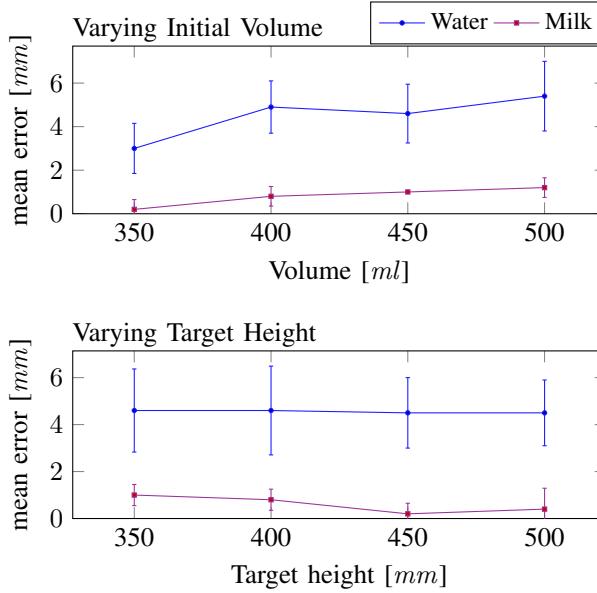
#### 4.3 Influence of Varying Target Heights

Here we look at the influence of the target height. We kept the initial volume constant at 400 ml and varied the target height to one of [30 mm, 40 mm, 50 mm, 60 mm]. For each height, we poured ten times for both water and milk. The blue bottle and blue cup were used each time for consistency.

The results can be seen in the bottom figure of Fig. 5. The mean errors and standard deviations remain fairly consistent for both liquids. So it does not appear as if different target heights have a large influence on the pouring error.

#### 4.4 Influence of Bottle Opening

In this experiment, we investigate the influence of the bottle opening. The silver bottle (referred to as the wide opening bottle) pictured on the left in Fig. 1 has an opening of 4.5 cm while the blue bottle (referred to as the small opening bottle), pictured on the right, has an opening of 2.5 cm. We conducted ten pours



**Fig. 5.** Mean error and standard deviation over ten pours for varying initial liquid volumes in the bottle and constant target height (top). Mean error and standard deviation over ten pours for varying target height with constant initial liquid volumes (bottom).

each for water and milk, using the wide opening bottle. The initial volumes and target heights were varied in the same manner as in Sec. 4.1 for each pour.

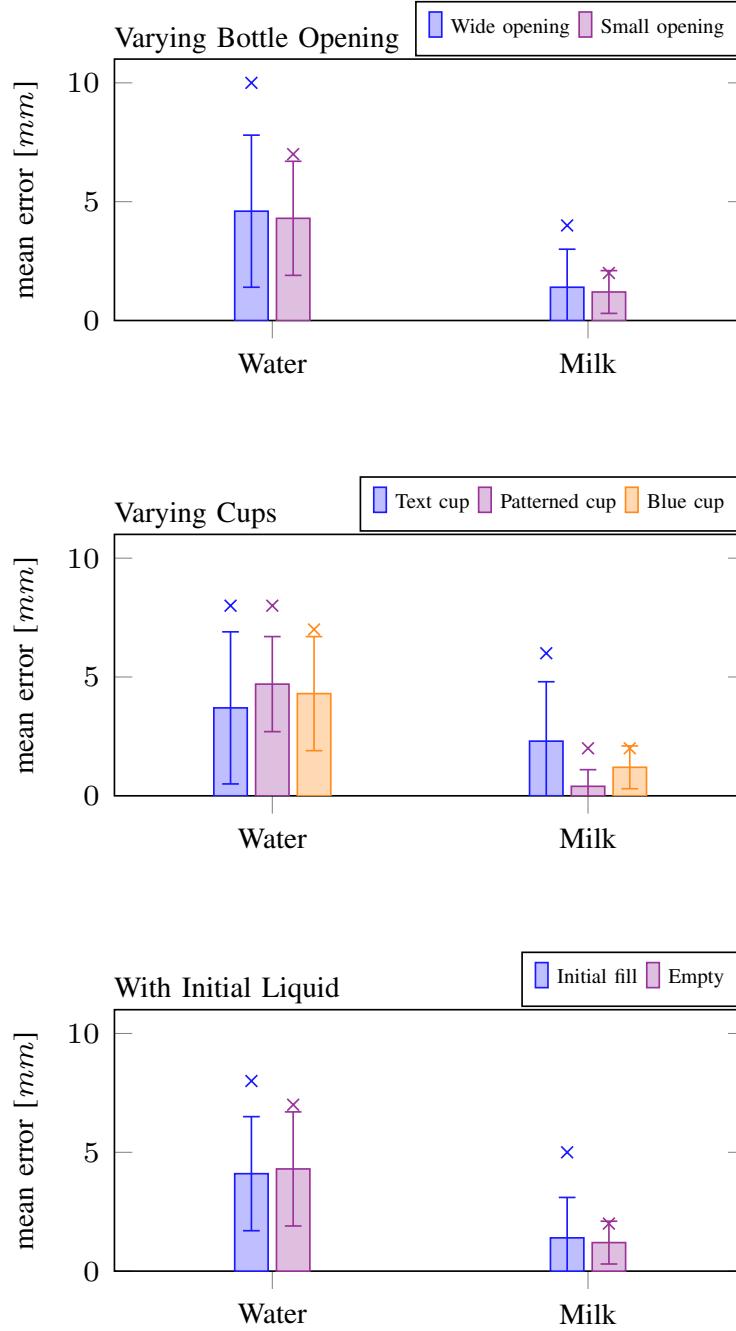
The results can be seen in the top figure of Fig. 6. We include the results for the small opening bottle from Sec. 4.1 for reference. A small increase in mean error (0.3 mm for water and 0.2 mm for milk) for the wide opening bottle can be seen. This can be expected since a wider opening makes it more difficult to control the flow of the fluid as it leaves the bottle. But overall, this experiment demonstrates that our system is able to deal with different openings.

#### 4.5 Influence of Different Cups

So far, each experiment was conducted with the blue cup, depicted on the right side in Fig. 1. In this section we investigate the influence of different cups, in particular whether the shape and cup bottom diameter affect the system. We refer to the leftmost cup in Fig. 1 as the *text cup*, the middle cup as the *patterned cup* and the rightmost cup as the *blue cup*. For both the text cup and the patterned cup, we performed ten pours each, using the blue bottle for water and milk. As in Sec. 4.1, we varied the initial volume of liquid in the bottle and the target height. We used the same approach for detecting the cup in each case.

The results can be seen in the middle figure of Fig. 6. For reference we include the results for the blue cup from Sec. 4.1. In general, the differences in mean

X



**Fig. 6.** Mean error and standard deviation over ten pouring experiments with a wide and a small opening bottle (top). Mean error and standard deviation over ten pours for each of the three different cups (middle). Mean error and standard deviation over ten pours for each liquid for pouring into a partly pre-filled cup (bottom). In all cases, the crosses represent the worst results over all ten pours.

error are minor between cups and there is no clear trend. The cup with the smallest bottom is the text cup, having a diameter of around 5 cm, compared to 6 cm for the patterned cup and 7.5 cm for the blue cup. In the case of water, where it is the refracted cup bottom that is being measured, one would expect a smaller cup bottom would be more problematic, but this does not appear to be the case. Furthermore, it was unclear beforehand whether the shape of the cups would affect Eq. (1) (due to reflections off the cup sides). But this does not appear to be an issue.

#### 4.6 Influence of Initial Liquid in Cup

In this experiment, we look at how the system can deal with initial liquid levels in the cup. For both water and milk, we performed ten pours each. In each pour, we chose an initial amount of liquid in the cup ranging between 10 –30 mm. We varied the initial volume in the bottle and ensured the target height was at least 10 mm more than the initial liquid amount. Once again, we used the blue bottle and blue cup.

The bottom figure in Fig. 6 shows the mean errors, standard deviations and maximum pour error. We also include the case of starting with an empty cup (see Sec. 4.1 for reference). The results show only minor differences between pouring into an empty cup and pouring into one with an initial amount of liquid. Accordingly our system is able to robustly deal with situations in which the cup is partly pre-filled.

#### 4.7 Comparison of Pouring Accuracy

As noted before, Schneck and Fox [13] report a mean error of 38 ml for water over 30 pours, using a system that also controls the rotation angle of the wrist joint gripping the bottle. For each of our cups, we take the worst 10-pour trial, which all occurred while pouring water, and converted the height errors to volume errors. This resulted in mean volume errors of 23.9 ml for the blue cup, 13.2 ml for the text cup and 30.5 ml for the patterned cup. The patterned cup is also the widest cup of the three, meaning errors in height result in larger volume errors. Overall, our system was able to achieve better accuracy in pouring, while under more extreme testing conditions.

### 5 Conclusion

In this paper, we presented a novel approach for liquid pouring that uses data received from an RGB-D camera and allows for pouring to a specific, user-defined height. Our approach tracks the liquid height during the pour which allows it to accurately stop as soon as the desired height has been reached. We conducted extensive experiments with our PR2 robot, which show that we are able to accurately pour both a transparent and an opaque liquid under various conditions. In future work we will look at reducing the dependency on user input

and improving the handling of delays between the perception and the controller so as to further improve the accuracy.

## References

1. Pan, Z., Manocha, D.: Motion planning for fluid manipulation using simplified dynamics. *Intelligent Robots and Systems (IROS)*, 4224–4231 (2016)
2. Tamosiunaite, M., Nemec, B., Ude, A., Wörgötter, F.: Learning to pour with a robot arm combining goal and shape learning for dynamic movement primitives. *Robotics and Autonomous Systems* 59, 910–922 (2011)
3. Langsfeld, J., Kaipa, K., Gentili, R., Reggia, J., Gupta, S.: Incorporating failure-to-success transitions in imitation learning for a dynamic pouring task. In *Workshop on Compliant Manipulation: Challenges and Control* (2014)
4. Rozo, L., Jiménez, P., Torras, C.: Force-based robot learning of pouring skills using parametric hidden Markov models. *Robot Motion and Control (RoMoCo)*, 224–232 (2013)
5. Elbrechter, C., Maycock, J., Haschke, R., Ritter, H.: Discriminating Liquids Using a Robotic Kitchen Assistant. *Intelligent Robots and Systems (IROS)*, 703–708 (2015)
6. Morris, N., Kutulakos, K.: Dynamic refraction stereo. *IEEE Transactions on Pattern Analysis and Machine Intelligence* 33, 1518–1531 (2011)
7. Mottaghi, R., Schenck, C., Fox, D., Farhadi, A.: See the Glass Half Full: Reasoning about Liquid Containers, their Volume and Content. *arXiv preprint arXiv:1701.02718* (2017)
8. Yamaguchi, A., Atkeson, C., Ogasawara, T.: Pouring skills with planning and learning modeled from human demonstrations. *International Journal of Humanoid Robotics* 12, 1550030 (2015)
9. Yamaguchi, A., Atkeson, C.: Stereo vision of liquid and particle flow for robot pouring. *Humanoid Robots (Humanoids)*, 1173–1180 (2016)
10. Yoshitaka, H., Fuhito, H., Takashi, T., Akihisa, O.: Detection of Liquids in Cups Based on the Refraction of Light with a Depth Camera Using Triangulation. *Intelligent Robots and Systems (IROS)*, 5049–5055 (2014)
11. Do, C., Schubert, T., Burgard, W.: A Probabilistic Approach to Liquid Level Detection in Cups Using an RGB-D Camera. *Intelligent Robots and Systems (IROS)*, 2075–2080 (2016)
12. Schenck, C., Fox, D.: Detection and tracking of liquids with fully convolutional networks. *arXiv preprint arXiv:1606.06266* (2016)
13. Schenck, C., Fox, D.: Visual closed-loop control for pouring liquids. *Robotics and Automation (ICRA)*, 2629–2636 (2017)
14. Fischler, M., Bolles, R.: Random sample consensus: A paradigm for model fitting with applications to image analysis and automated cartography. *Communications of the ACM* 24, 381–395 (1981)
15. Andersen, M.R., Jensen, T., Lisouski, P., Mortensen, A.K., Hansen, M.K., Gregersen, T., Ahrendt, P.: Kinect Depth Sensor Evaluation for Computer Vision Applications. *Thechnical Report ECE-TR-6* (2012)